

CS2110 Fall 2015 Assignment A1 PhD Genealogy

Introduction

Website <http://genealogy.math.ndsu.nodak.edu> contains the PhD genealogy of about 189,139 mathematicians and computer scientists, showing their PhD advisors and advisees. Foster's intellectual great grandfather is Prof. Bob Constable of our CS Department (and was the first Dean of CIS). Both Foster and Gries can trace their intellectual ancestry back to Gottfried Wilhelm Leibniz (1646–1716), who dreamed of a “general method in which all truths of the reason would be reduced to a kind of calculation”. Leibniz foresaw symbol manipulation as we know it today.

It's a laborious, error-prone task to search the genealogy website by hand and construct a tree of someone's PhD ancestors, so we wrote a Java program to construct the tree for a given person. It uses a class `PhD` much like the one you will build, but it has many more fields because of the complexity of the information on that website. The program also uses a class that allows one to read a web page. At the appropriate time, we can show you this program and discuss its construction, so you can learn how to write programs that crawl webpages.

The intellectual genealogies of Foster and Gries can be found on the assignments page of the course website. Gries has 256 distinct ancestor's going back to the year 1315 and including Copernicus. Foster has 96.

Your task in this assignment is to develop a Java class `PhD` that will maintain information about the PhD of a person, giving the date and advisor(s), and a JUnit class `PhDTester` to maintain a suite of test cases for class `PhD`. Note that the term PhD is not used in all countries! Foster has a PhD, but Gries's degree is a *Dr. Rerum Natura* from MIT (Munich Institute of Technology). It is abbreviated *Dr. red nat*, which Gries speaks as *rare nut*. In this assignment, we use only the term PhD.

Your last task before submitting the assignment will be to tell us how much time you spent on this assignment, so please keep track. This will allow us to publish the mean, median, and maximum times, so you have an idea how you are doing relative to others. It also helps us ensure that we don't require too much of your time in this course.

Learning objectives

- Gain familiarity with the structure of a class within a record-keeping scenario (a common type of application)
- Learn about and practice reading carefully.
- Work with examples of good Javadoc specifications to serve as models for your later code.
- Learn the code presentation conventions for this course (Javadoc specifications, indentation, short lines, etc.), which help make your programs readable and understandable.
- Learn and practice incremental coding, a sound programming methodology that interleaves coding and testing.
- Learn about and practice thorough testing of a program using JUnit testing.
- Learn to write *class invariants*.
- Learn about *preconditions* of a method (requirements of a call on the method that the caller must follow) and the use of the Java assert statement for checking preconditions.

The methods to be written are short and simple; the emphasis in this assignment is on “good practices”, not complicated computations.

Reading carefully

At the end of this document is a checklist of items for you to consider before submitting A1, showing how many points each item is worth. Check each item *carefully*. A low grade is almost always due to lack of attention to detail and to not following instructions —not to difficulty understanding OO. At this point, we ask you to visit the webpage on the website of Fernando Pereira, research director for Google:

<http://earningmyturns.blogspot.com/2010/12/reading-for-programmers.html>

Did you read that webpage carefully? If not, read it now! The best thing you can do for yourself —and us— at this point is to read carefully. This handout contains many details. Save yourself and us a lot of anguish by reading carefully as you do this assignment.

Collaboration policy

You may do this assignment with one other person. If you are going to work together, then, as soon as possible —and certainly before you submit the assignment— get on the CMS for the course and do what is required to form a group. Both people must do something before the group is formed: one proposes, the other accepts. If you need help with the CMS, visit www.cs.cornell.edu/Projects/CMS/userdoc/.

If you do this assignment with another person, you must *work together*. It is against the rules for one person to do some programming on this assignment without the other person sitting nearby and helping. You should take turns “driving” —using the keyboard and mouse. If you are the weaker of two students on a team and you let your teammate do more than their share, you are hurting only yourself. You can’t learn without doing.

With the exception of your CMS-registered partner, you may not look at anyone else’s code, in any form, or show your code to anyone else, in any form. You may not show or give your code to another person in the class. While you can talk to others, your discussions should not include writing code and copying it down.

Getting help

If you don’t know where to start, if you don’t understand testing, if you are lost, etc., please SEE SOMEONE IMMEDIATELY —a course instructor, a TA, a consultant. Do not wait. A little in-person help can do wonders.

Using the Java assert statement to test preconditions

A *precondition* is a constraint on the parameters of a method, and it is up to the user to ensure that method calls satisfy the precondition. If a call does not, the method can do whatever it wants.

However, especially during testing and debugging, it is useful to use Java assert statements at the beginning of a method to check that the precondition is true. For example, if the precondition is “this person’s name is at least one character long”, one can use an assert statement like the following (using the `name` for the field):

```
assert name != null  &&  name.length() >= 1;
```

The additional test `name != null` is there to protect against a null-pointer exception, which will happen if the argument corresponding to `name` in the call is `null`. [This is important!]

In this assignment, all preconditions of methods must be checked using assert statements in the method. Write the assert statements as the first step in writing the method body, so that they are always there during testing and debugging. Also, when you generate a new JUnit class, make sure the VM argument `-ea` is present in the Run Configuration. You will find assert statements helpful in testing and debugging.

How to do this assignment

Scan the whole assignment before starting. Then, develop class `PHD` and test it using class `PHDTester` in the following incremental, sound way. This methodology will help you complete this (and other) programming tasks quickly and efficiently. If we detect that you did not develop it this way, points will be deducted.

1. In Eclipse, create a new project, called `a1Genealogy`. In `a1Genealogy`, create a new class, called `PhD`. It should *not* be in a package, and it does not need a method `main` (but you may add one if you want). Insert the following as the first lines of file `PhD.java`:

```
/** NetId: nnnnn, nnnnn. Time spent: hh hours, mm minutes.
    An instance maintains info about the PhD of a person. */
```

Remove the constructor with no parameters, since it will not be used and its use can leave an object in an inconsistent state (see below, the class invariant).

2. In class `PhD`, declare the following fields (you choose the names of the fields), which are meant to hold information describing a person with a PhD. Make these fields private and properly comment them (see the "[The class invariant](#)" section below).

Note: you must break long lines (including comments) into two or more lines so that the reader does not have to scroll right to read them. This makes your code much easier for us and *you* to read. A good guideline: No line is more than 80 characters long.

- ▶ `name` (a `String`). Name of the person with a PhD, a `String` of length > 0 .
- ▶ `year` PhD was awarded (an `int`).
- ▶ `month` PhD was awarded (an `int`). In range 1..12 with 1 being January, etc.
- ▶ `gender` of the person (a `char`). 'M' means male and 'F' means female.
- ▶ `first advisor` of this person (of type `PhD`). The first PhD advisor of this person —null if unknown.
- ▶ `second advisor` of this person (of type `PhD`). The second advisor of this person —null if unknown or if the person had only one advisor. If the first-advisor field is null, the second advisor field must be null.
- ▶ `number` of PhD advisees of this person.

About the field that contains the number of advisees: The user *never* gives a value for this field; it is completely under control of the program. For example, whenever a `PhD p` is given an advisor `m`, `m`'s number of advisees must be increased by 1. *It is up to the program, not the user, to increase the field.*

The class invariant. Recall that comments should accompany the declarations of all fields to describe what each field means, what constraints hold on the fields, and what the legal values are for the fields. For example, for the `name-of-the person` field, state in a comment that the field contains the person's name and must be a string of at least 1 character. The collection of the comments on these fields is called the *class invariant*. Here is an example of a declaration with a suitable comment. Note that the comment does *not* give the type (since it is obvious from the declaration), it does not use noise phrases like "this field contains ...", and it *does* contain constraints on the field.

```
int month; // month the PhD was awarded. In range 1..12, with 1 meaning January, etc.
```

Note again that we did not put "(an int)" in the comment. That information is already known from the declaration. Don't put such unnecessary things in the comments.

Whenever you write a method (see below), look through the class invariant and convince yourself that the class invariant still holds when the method terminates. This habit will help you prevent or catch bugs later on.

3. In Eclipse, start a new JUnit test class and call it `PhDTester`. You can do this using menu item **File —> New —> JUnit Test Case** (add the JUnit 4 library, if asked).
4. Below, we describe four *groups* A, B, C, and D of methods. Work with *one* group at a time, performing steps (1)..(4). **Do not go on to the next group until the group you are working on is thoroughly tested and correct.**
 - (1) Write the Javadoc specifications for each method in that part. Make sure they are complete and correct — look at the specs we give you below. Copy-and-paste from this handout makes this easy.
 - (2) Write the method bodies, starting with assert statements for the preconditions.
 - (3) Write *one* test procedure for this group in class `PhDTester` and add test cases to it for all the methods in the group.
 - (4) Test the methods in the group thoroughly.

Discussion of the groups of methods. The descriptions below represent the level of completeness and precision required in Javadoc specification-comments. In fact, you may (should) copy and paste these descriptions to create the first draft of your Javadoc comments. If you do not cut and paste, adhere to the conventions we use, such as using the prefix “Constructor: ...” and double-quotes to enclose an English boolean assertion. Using a consistent set of good conventions in this class will help us all.

Method specifications do not mention fields because the user may not know what the fields are, or even if there are fields. The fields are private. Consider class `JFrame`: you know what methods it has but not what fields, and the method specifications do not mention fields. In the same way, a user of your class `PhD` will know the methods but not the fields.

The names of your methods must match those listed below exactly, including capitalization. The number of parameters and their order must also match: any mismatch will cause our testing programs to fail and will result in loss of points for correctness. Parameter names will not be tested —change the parameter names if you want.

In this assignment, you may *not* use if-statements, conditional expressions, or loops.

Group A: The first constructor and the getter methods of class `PhD`.

Constructor	Description (and suggested javadoc specification)	
<code>PhD(String n, char g, int y, int m)</code>	Constructor: an instance for a person with name <code>n</code> , gender <code>g</code> , PhD year <code>y</code> , and PhD month <code>m</code> . Its advisors are unknown, and it has no advisees. Precondition: <code>n</code> has at least 1 character, <code>m</code> is in 1..12, and <code>g</code> is 'M' for male or 'F' for female.	
Getter Method	Description (and suggested javadoc specification)	Return Type
<code>getName()</code>	Return this person's name.	<code>String</code>
<code>getYear()</code>	Return the year this person got their PhD.	<code>int</code>
<code>getMonth()</code>	Return the month this person got their PhD.	<code>int</code>
<code>isMale()</code>	Return the value of the sentence "this person is a male."	<code>boolean</code>
<code>getFirstAdvisor()</code>	Return this PhD's first advisor (null if unknown).	<code>PhD</code> (not <code>String</code> !)
<code>getSecondAdvisor()</code>	Return this PhD's second advisor (null if unknown or non-existent).	<code>PhD</code> (not <code>String</code> !)
<code>numAdvisees()</code>	Return the number of PhD advisees of this person.	<code>int</code>

Consider the constructor. Based on its specification, figure out what value it should place in each of the 7 fields to make the class invariant true. Then, write a procedure named `testConstructor1` in `PhDTester` to make sure that the constructor fills in ALL fields correctly. The procedure should: Create one `PhD` object using the constructor and then check, using the getter methods, that all fields have the correct values. Since there are 7 fields, there should be 7 `assertEquals` statements. As a by-product, all getter methods are also tested.

We advise creating a second `PhD` (in `testConstructor1`) of the other sex than the one first created and testing —using function `isMale()`— that its sex was properly stored.

Group B: the setter/mutator methods. Note that methods `addFirstAdvisor` and `addSecondAdvisor` may have to change fields of both this `PhD` and its parent in order to maintain the class invariant —the advisor's number of advisees changes!

When testing the setter methods, you will have to create one or more `PhD` objects, call the setter methods, and then use the getter methods to test whether the setter methods set the fields correctly. Good thing you already tested the getters! Note that these setter methods may change more than one field; your testing procedure should check that *all* fields that may be changed are changed correctly.

We are *not* asking you to write methods that change an existing father or mother to a different PhD. This would require if-statements, which are not allowed. Read preconditions of methods carefully.

Setter Method	Description (and suggested javadoc specification)
addFirstAdvisor(PhD p)	Add p as this person's first PhD advisor. Precondition: this person's first advisor is unknown and p is not null.
addSecondAdvisor(PhD p)	Add p as this person's second advisor. Precondition: This person's first advisor is known, their second advisor is unknown, p is not null, and p is different from this person's first advisor.

Group C: Two more constructors. The test procedure for group C has to create a PhD using the constructor given below. This will require first creating two PhD objects using the first constructor and then checking that the new constructor sets *all* 7 fields properly —and also the number of advisees of ad1 and ad2.

Constructors	Description (and suggested javadoc specification)
PhD(String n, char g, int y, int m, PhD adv)	Constructor: a PhD with name n, gender g, PhD year y, PhD month m, first advisor adv, and no second advisor. Precondition: n has at least 1 char, g is 'F' for female or 'M' for male, m is in 1..12, and adv is not null.
PhD(String n, char g, int y, int m, PhD adv1, PhD adv2)	Constructor: a PhD with name n, gender g, PhD year y, PhD month m, first advisor adv1, and second advisor adv2. Precondition: n has at least 1 char, g is 'F' for female or 'M' for male, m is in 1..12, adv1 and adv2 are not null, and adv1 and adv2 are different.

Group D: Write two comparison methods —to see which of two people got their PhD first and to see whether two people are “intellectual siblings” (i.e. they are not the same object and they have a non-null advisor in common). Write these using only boolean expressions (with !, &&, and || and relations <, <=, ==, etc.). *Do not use if-statements, switches, addition, multiplication, etc.* Each is best written as a single return statement.

Comparison Method	Description (and suggested javadoc specification)	Return type
isYoungerThan(PhD p)	Return value of "this person got their PhD after p did." Precondition: p is not null.	boolean
isPhDSibling(PhD p)	Return value of "this person and p are intellectual siblings." Note: if p is null, they are not siblings.	boolean

- In Eclipse, click menu item **Project -> Generate Javadoc**. In the window that opens, make sure you are generating Javadoc for project, a1PhD, using visibility **public**, and storing it in a1PhD/doc. Then open doc/index.html. You should see your method and class specifications. Read through them from the perspective of someone who has not read your code. Fix the comments in class PhD, if necessary, so that they are appropriate to that perspective. You *must* be able to understand everything there is to know about writing a call on each method from the specification that you see by clicking the Javadoc button —that is, without knowing anything about the private fields. Thus, the fields should not be mentioned.

Then, and only then, add a comment at the top of file PhD.java saying that you checked the Javadoc output and it was OK.

- Check carefully that each method that adds an advisor for a PhD updates the advisor's number of advisees. Four methods do this.

7. Review the learning objectives and reread this document to make sure your code conforms to our instructions. Check each of the following, one by one, carefully. Note that 50 points are given for the items below and 50 points are given for actual correctness of methods.
 - 5 Points. Are all lines short enough that horizontal scrolling is not necessary (about 80 chars is long enough). Do you have a blank line before the specification of each method and no blank line after it?
 - 10 Points. Is your class invariant correct —are all fields defined and all field constraints given?
 - 5 Points. Is the name of each method and the types of its parameters exactly as stated in step 4 above? (The simplest way to do this was to copy and paste.) More points may be deducted if we have difficulty fixing your submission so that it compiles with our grading program.
 - 10 Points. Are all specifications complete, with any necessary preconditions? Remember, we specified every method carefully and suggested copying our specs and pasting them into your code. Are they in Javadoc comments?
 - 5 points. Do you have assert statements in each method that has a precondition to check that precondition?
 - 5 Points. Did you check the Javadoc output and then put a comment at the top of class PhD?
 - 10 Points. Did you write *one* (and only one) testing method for each of the groups A, B, C, and D of step 4? Thus, do you have four (4) test procedures? Does each procedure have a name that gives the reader an idea what the procedure is testing, so that a specification is not necessary? Did you properly test? For example, in testing each constructor, did you make sure to test that all 7 fields have the correct value? Do you have enough test cases? For example, testing whether one date comes before another date, when each is given by a month and a year, probably requires at least 5 test cases.
8. Change the first line of file PhD.java: replace “nnnnn” by your netids, “hh” by the number of hours spent, and “mm” by the number of minutes spent. If you are doing the assignment alone, remove the second “nnnn”. For example, suppose foster and gries spent 4 hours and 30 minutes. Then the first line would be as shown below.

```
/** NetIds: jnf27, djg17. Time spent: 4 hours, 30 minutes.
```

Being careful in changing this line will make it easier for us to automate the process of calculating the median, mean, and max times. Help us out and be careful.

9. Upload files PhD.java and PhDTester.java on the CMS by the due date. Do not submit any files with the extension/suffix .java~ (with the tilde) or .class. It will help to set the preferences in your operating system so that extensions always appear.

